


Hands on Deep Learning



VVV17 Winter School on Humanoid Robot Programming
Santa Margherita Ligure
Feb. 7 2017


Outline

1. The iCubWorld Project

2. Deep Learning using Caffe

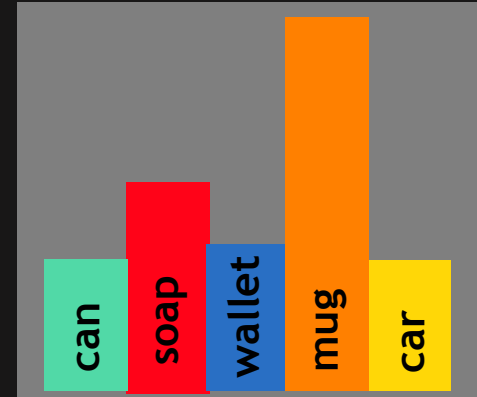

Teaching iCub to See: HRI Framework

Verbal
instructions
of a “teacher”



iCub learning
the object

On the Fly Object Recognition



Disparity
Segmentation

Motor
Control



Representation
Extraction

learn <obj>
forget <obj>
what is this?

Speech

Interaction
Manager


Classifier

The Need for an iCubWorld Dataset

What is the Visual World of a Robot?

- not so many object/instances

<http://image-net.org/>



The Need for an iCubWorld Dataset

What is the Visual World of a Robot?

- not so many object/instances
- lots of viewpoint changes



The Need for an iCubWorld Dataset

What is the Visual World of a Robot?

- not so many object/instances
- lots of viewpoint changes
- ...
- uninformative background
- self supervised
- ...

→ Different from usual vision tasks!!

iCubWorld Datasets

<https://robotology.github.io/iCubWorld/>



2012-13
1st and 2nd iCW Releases

2014-15
iCubWorld28

2016
iCubWorld - Transformations

Latest Release: iCubWorld - Transformations



20 categories

10 object instances per category



Latest Release: iCubWorld - Transformations

2D ROT:



3D ROT:



SCALE:



BKG:



MIX:



20 categories

10 object instances per category

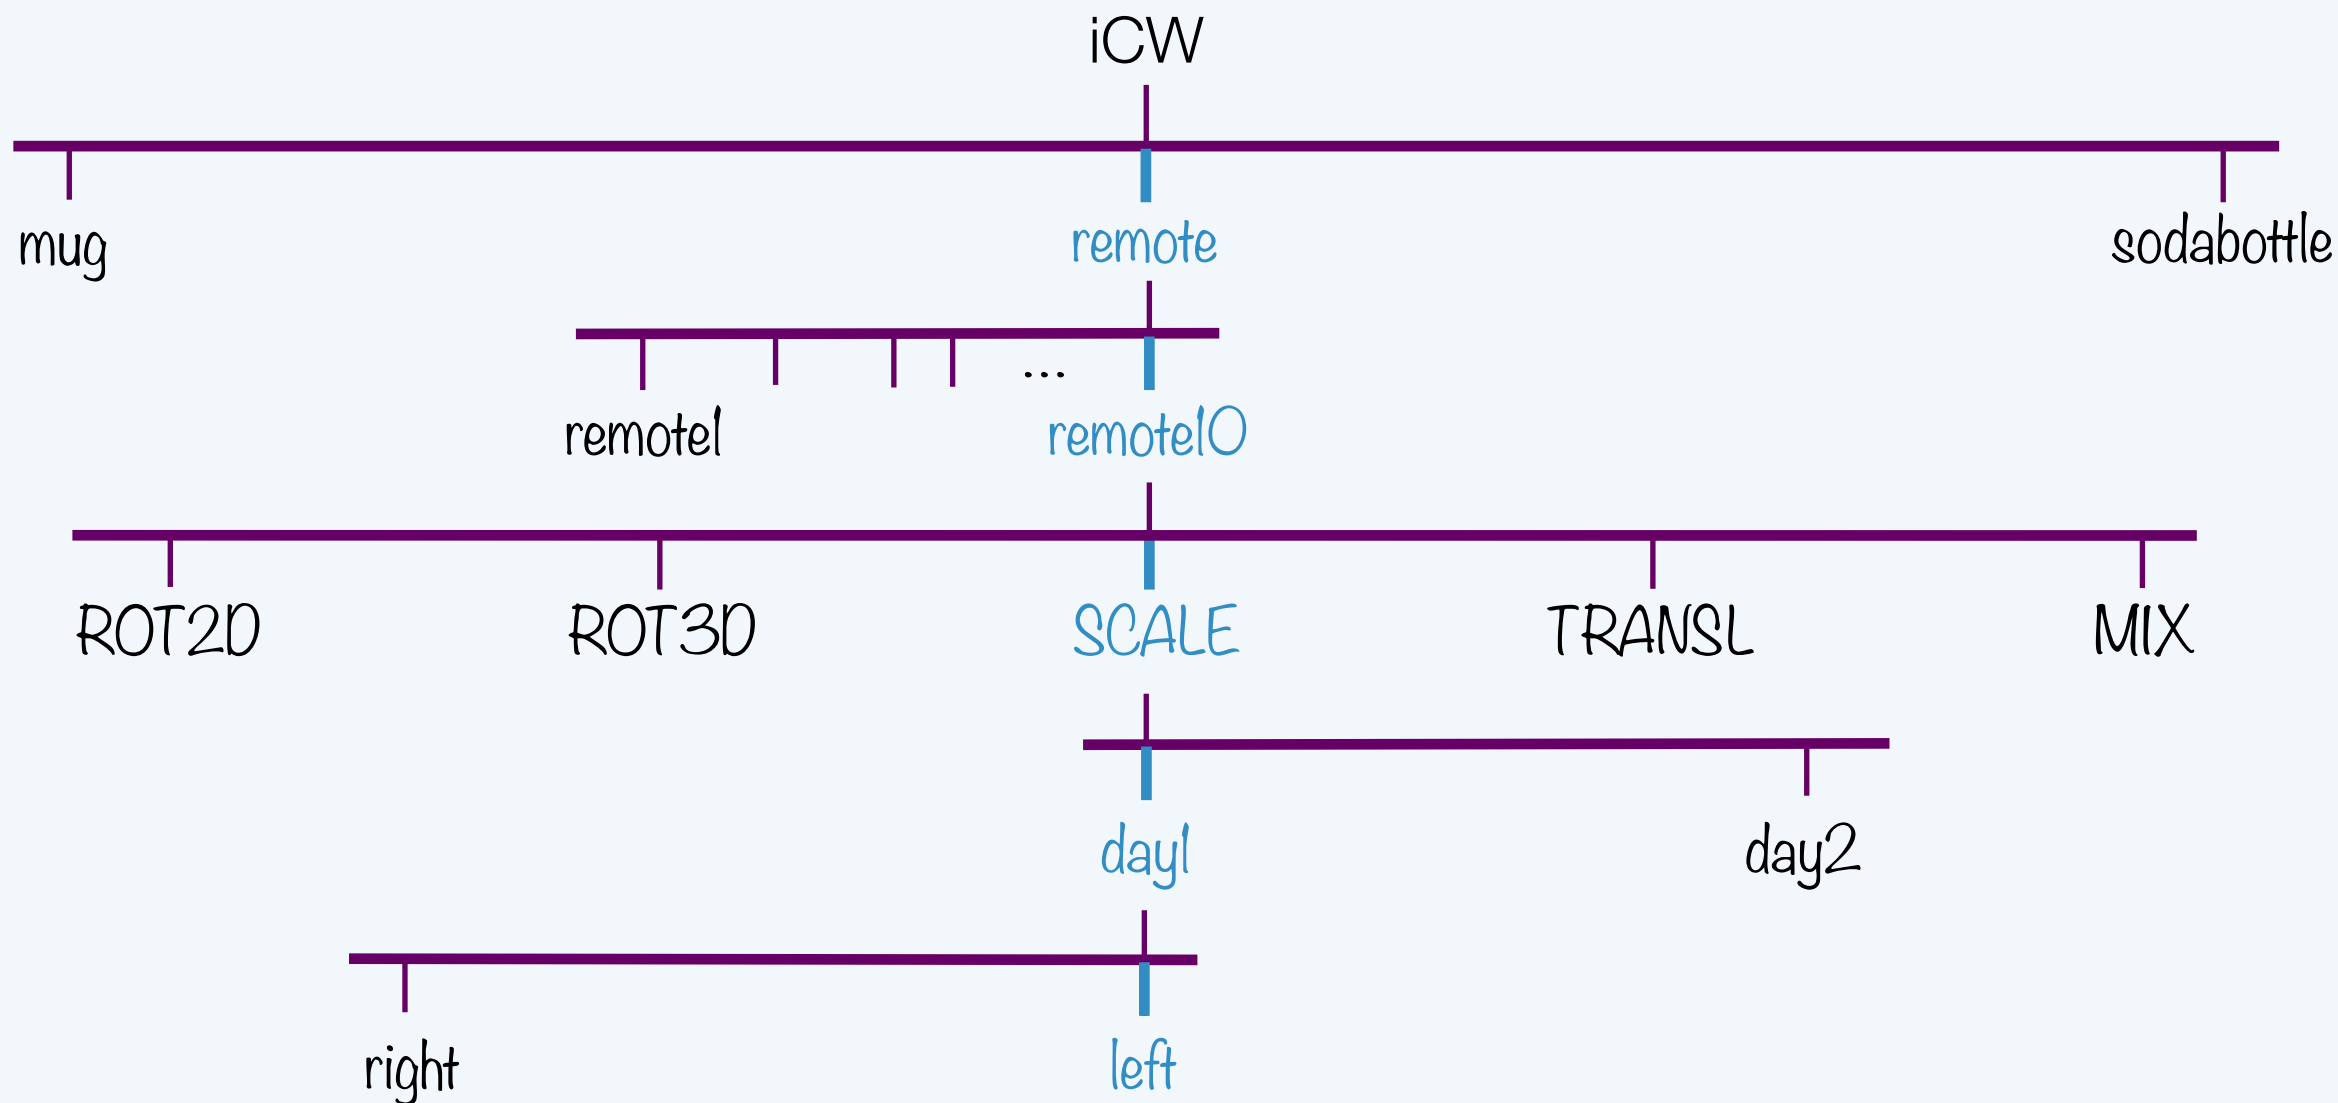
10 sequences per object:

5 transformations

2 days

stereo sequences

iCubWorld (iCW): subset for the labs



$3 \times 10 \times 5 \times 2 \times 1 \times (\sim 150 \text{ frames}) \rightarrow \sim 50\text{k images}$

Outline

1. The iCubWorld Project

2. Deep Learning using Caffe

Caffe

(Convolutional Architecture for Fast Feature Embedding)

C++ framework

command-line, **Python** and **Matlab** interfaces

to define and train **(Convolutional) Neural Networks**
(on CPUs/CUDA GPUs)

Web: <http://caffe.berkeleyvision.org/>

GitHub: <https://github.com/BVLC/caffe>

Community: [Caffe Users - Google Groups](#)

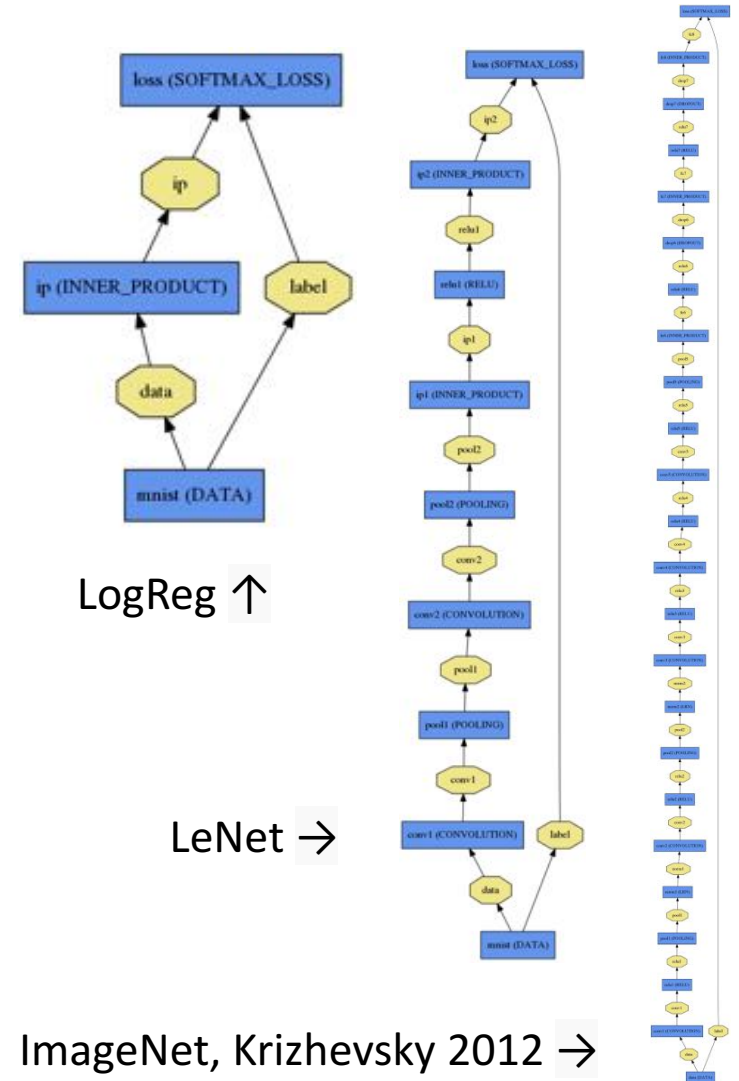
Caffe

(Convolutional Architecture for Fast Feature Embedding)


Nets

Layers

Blobs



Net



A network is a set of **Layers** and their connections.

Data and derivatives flow through the net as **Blobs** (an array interface).

Forward/Backward are the essential **Net** computations.


Layers

Defined by **setup**, **forward** and **backward** computations:

setup: run once for initialization


forward: make output given input

backward: make gradient of output
- w.r.t. bottom
- w.r.t. parameters




The Net's **Forward** and **Backward** passes are composed of the layers' steps (*Compositional Modeling*)

LAYER CATALOGUE (on Caffe website)




- ✓ "Convolution"
- ✓ "Pooling"
- ✓ "LRN" (normalization)
- ✓ "InnerProduct"
- ✓ "ReLU" (non linearity)
- ✓ ...
- ✓ **DATA layers** (include image pre-processing steps):
 - ✓ database ("Data")
 - ✓ in-memory ("MemoryData")
 - ✓ image list ("ImageData")
- ✓ **LOSS layers**
 - ✓ softmax ("SoftmaxWithLoss")
 - ✓ accuracy ("Accuracy")

(Common) Image Preprocessing Steps



- ✓ “Convolution”
- ✓ “Pooling”
- ✓ “LRN” (normalization)
- ✓ “InnerProduct”
- ✓ “ReLU” (non linearity)
- ✓ ...
- ✓ **DATA layers** (include image pre-processing steps):
 1. Subtract mean image of the training set (or mean pixel)
 2. Crop of expected size (e.g. 227x227x3)
- ✓ **LOSS layers**
 - ✓ softmax (“SoftmaxWithLoss”)
 - ✓ accuracy (“Accuracy”)

Blobs




4-dimensional array stored in a C-contiguous fashion:

number N x channel C x height H x width W


Hides the overhead of CPU/GPU computation by synchronizing as needed and allocating memory on host/device lazily.

Provides a unified memory interface:



Data

Number x Channel x Height x Width
256 x 3 x 227 x 227 for CaffeNet train input



Parameter: Convolution Weight

N Output x K Input x Height x Width
96 x 3 x 11 x 11 for CaffeNet conv1



Parameter: Convolution Bias

96 x 1 x 1 x 1 for CaffeNet conv1

Solver (Optimizer)

Initializes the model and the training.

At each iteration:

forward → output and loss

backward → gradients

parameter updates

solver state update

Periodic validation.

Periodic snapshots of the model and solver state.

Many different kinds implemented in Caffe, e.g.:

- SGD
- Nesterov Accelerated Gradient
- Adam
- ...

Hands On

1. Warm Up

- running out-of-the-box
- opening the box

2. Assignment

- change configuration files and run different fine-tuning protocols
- [optional] suggestions for deeper exploration

Hands On

1. Warm Up

- running out-of-the-box
- opening the box

2. Assignment

- change configuration files and run different fine-tuning protocols
- [optional] suggestions for deeper exploration

Running out-of-the-box

Follow instructions here: https://github.com/vvv-school/tutorial_dl-tuning

1. run the tester `train_and_test_net_tester.sh`
(you should already have done this)
2. run the example `train_and_test_net.sh`

Opening the box

Look at the content of the example:


```
$ cd $LAB_DIR/tutorial_dl-tuning/id_2objects_caffenet
```

- all-0/train_val.prototxt
- all-0/deploy.prototxt
- all-0/solver.prototxt
- images_lists/train.txt
- images_lists/val.txt
- images_lists/test.txt
- images_lists/labels.txt

train_and_test_net.sh

What's happening?

Fine-tuning CaffeNet on a 2-class identification task



train_val.prototxt

batch_size (train)

...

lr_mult (conv1-5)

lr_mult (fc6)

lr_mult (fc7)

lr_mult (fc8N)

...

[num_output]

solver.prototxt

max_iter

[test_interval,
display]

base_lr



train_val.prototxt

batch_size (train)

...

lr_mult (conv1-5)

lr_mult (fc6)

lr_mult (fc7)

lr_mult (fc8N)

...

[num_output]

solver.prototxt

max_iter

[test_interval,
display]

base_lr

1. which layer fine-tune/learn from scratch
2. with which learning rate(s)
3. varying batch size

train_val.prototxt

```
batch_size (train)  
batch_size (validation)
```

...

```
lr_mult (conv1-5)  
lr_mult (fc6)  
lr_mult (fc7)  
lr_mult (fc8N)
```

...

```
num_output
```


deploy.prototxt

```
fc6
```

```
fc7
```

```
fc8N
```

```
num_output
```



train_and_test_net.sh

First defines paths (to caffe executables, to images, etc.), then:

1. create databases (train & val LMDBs and mean.binaryproto):

```
${Caffe_ROOT}/build/tools/convert_imageset \  
    --resize_width=256 --resize_height=256 --shuffle \  
    ${IMAGES_DIR} ${FILELIST_TRAIN} ${LMDB_TRAIN}  
→ lmdb_train (lmdb_val)
```

```
${Caffe_ROOT}/build/tools/compute_image_mean \  
    ${LMDB_TRAIN} ${BINARYPROTO_MEAN}  
→ mean.binaryproto
```

train_and_test_net.sh

2. train:

```
${Caffe_ROOT}/build/tools/caffe \  
    train -solver ${SOLVER_FILE} -weights ${WEIGHTS_FILE} \  
    --log_dir=${TUTORIAL_DIR}/${EX}/${PROTOCOL}  
→ icw_iter_*.caffemodel, icw_iter_*.solverstate, caffe.INFO
```

3. parse caffe.INFO:

```
${TUTORIAL_DIR}/scripts/parse_caffe_log.sh \  
    ${TUTORIAL_DIR}/${EX}/${PROTOCOL}/caffe.INFO  
→ caffeINFOtrain.txt, caffeINFOval.txt (tables)
```

train_and_test_net.sh

4. plot caffeINFOtrain.txt, caffeINFOval.txt (e.g. with MATLAB):

```
matlab -nodisplay -nodesktop -r \  
    "addpath('${TUTORIAL_DIR}/scripts');  
    try  
        plot_log('${TUTORIAL_DIR}/${EX}/${PROTOCOL}');  
    catch;  
    end;  
    quit"
```

→ matlab_caffeINFO_acc.png, matlab_caffeINFO_loss.png, matlab_caffeINFO_lr.png

train_and_test_net.sh

5. eventually choose best (or last) epoch:

```
snap_list=(`ls -t icw_iter*.caffemodel`)  
FINAL_SNAP=${TUTORIAL_DIR}/${EX}/${PROTOCOL}/${snap_list[0]}  
FINAL_MODEL=${TUTORIAL_DIR}/${EX}/${PROTOCOL}/final.caffemodel  
mv ${FINAL_SNAP} ${FINAL_MODEL}  
rm ${TUTORIAL_DIR}/${EX}/${PROTOCOL}/icw_iter_*.solverstate  
rm ${TUTORIAL_DIR}/${EX}/${PROTOCOL}/icw_iter_*.caffemodel
```

3. test final.caffemodel:

```
${TUTORIAL_DIR}/scripts/src/build/classify_image_list_vvv17 \  
    ${DEPLOY_FILE} ${FINAL_MODEL} ${BINARYPROTO_MEAN} \  
    ${LABELS_FILE} ${IMAGES_DIR} ${FILELIST_TEST} \  
    ${TUTORIAL_DIR}/${EX}/${PROTOCOL}/test_acc_day1.txt  
    ${PRINT_PREDICTIONS} ${IMG_DELAY}
```

→ test_acc_day1.txt

Hands On

1. Warm Up

- running out-of-the-box
- opening the box


2. Assignment

- change configuration files and run different fine-tuning protocols
- [optional] suggestions for deeper exploration

Assignment

1. Apply the protocol we have used in the tutorial on a different task
2. Try fine-tuning also fc7, fc6
3. Try fine-tuning all layers
4. Try learning fc7, fc6 from scratch
5. Vary the batch size

name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
all-0	32	lr: 0	lr: 0	lr: 0	lr: 1e-2 from scratch	6	?	?



name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
all-0	32	lr: 0	lr: 0	lr: 0	lr: 1e-2 from scratch	6	?	?
conv-0_fc6-2_fc7-2	32	lr: 0	lr: 1e-2	lr: 1e-2	lr: 1e-2 from scratch	6	?	?
conv-0_fc6-3_fc7-3	32	lr: 0	lr: 1e-3	lr: 1e-3	lr: 1e-2 from scratch	6	?	?
all-3	32	lr: 1e-3	lr: 1e-3	lr: 1e-3	lr: 1e-2 from scratch	6	?	?
conv-0_fc6N-2_fc7N-2	32	lr: 0	lr: 1e-2 from scratch	lr: 1e-2 from scratch	lr: 1e-2 from scratch	18	?	?
conv-0_fc6N-4_fc7N-4	32	lr: 0	lr: 1e-4 from scratch	lr: 1e-4 from scratch	lr: 1e-2 from scratch	36	?	?
all-3_batch-8	8	lr: 1e-3	lr: 1e-3	lr: 1e-3	lr: 1e-2 from scratch	6	?	?
all-0_batch-1	1	lr: 0	lr: 0	lr: 0	lr: 1e-2 from scratch	6	?	?

name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
all-0	32	lr: 0	lr: 0	lr: 0	lr: 1e-2 from scratch	6	?	?
conv-0_fc6-2_fc7-2	32	lr: 0	lr: 1e-2	lr: 1e-2	lr: 1e-2	6	?	?
conv-0_fc6-3							?	?
all-3							?	?
conv-0_fc6N-2							?	?
conv-0_fc6N-4							?	?
all-3_batch-8	8	lr: 1e-3	lr: 1e-3	lr: 1e-3	lr: 1e-2 from scratch	6	?	?
all-0_batch-1	1	lr: 0	lr: 0	lr: 0	lr: 1e-2 from scratch	6	?	?

Send your results to:
<https://goo.gl/forms/3kn1jV1jcQizlBa92>

Things you may want to try...

Try changing the task reducing the number of example frames per object in the training set:

1. create a new exercise folder (call it e.g. `id_10objects_less_frames`) by copying `id_10objects_caffenet`
2. open `id_10objects.m` in the new folder and change `setlist.max_frames.train` (try e.g. 20, or 10, or 5)
3. open `example_task_generation.m` and modify the paths to
 - (i) use the configuration file at point 2)
 - (ii) create training, validation and test sets in the new folder→ run the script to generate the new training set
4. implement the protocols `all-0` and `all-3` on this new task and run them
→ how much does the recognition accuracy decrease? is there a better protocol?

Things you may want to try...

Try to perform an object categorization task:

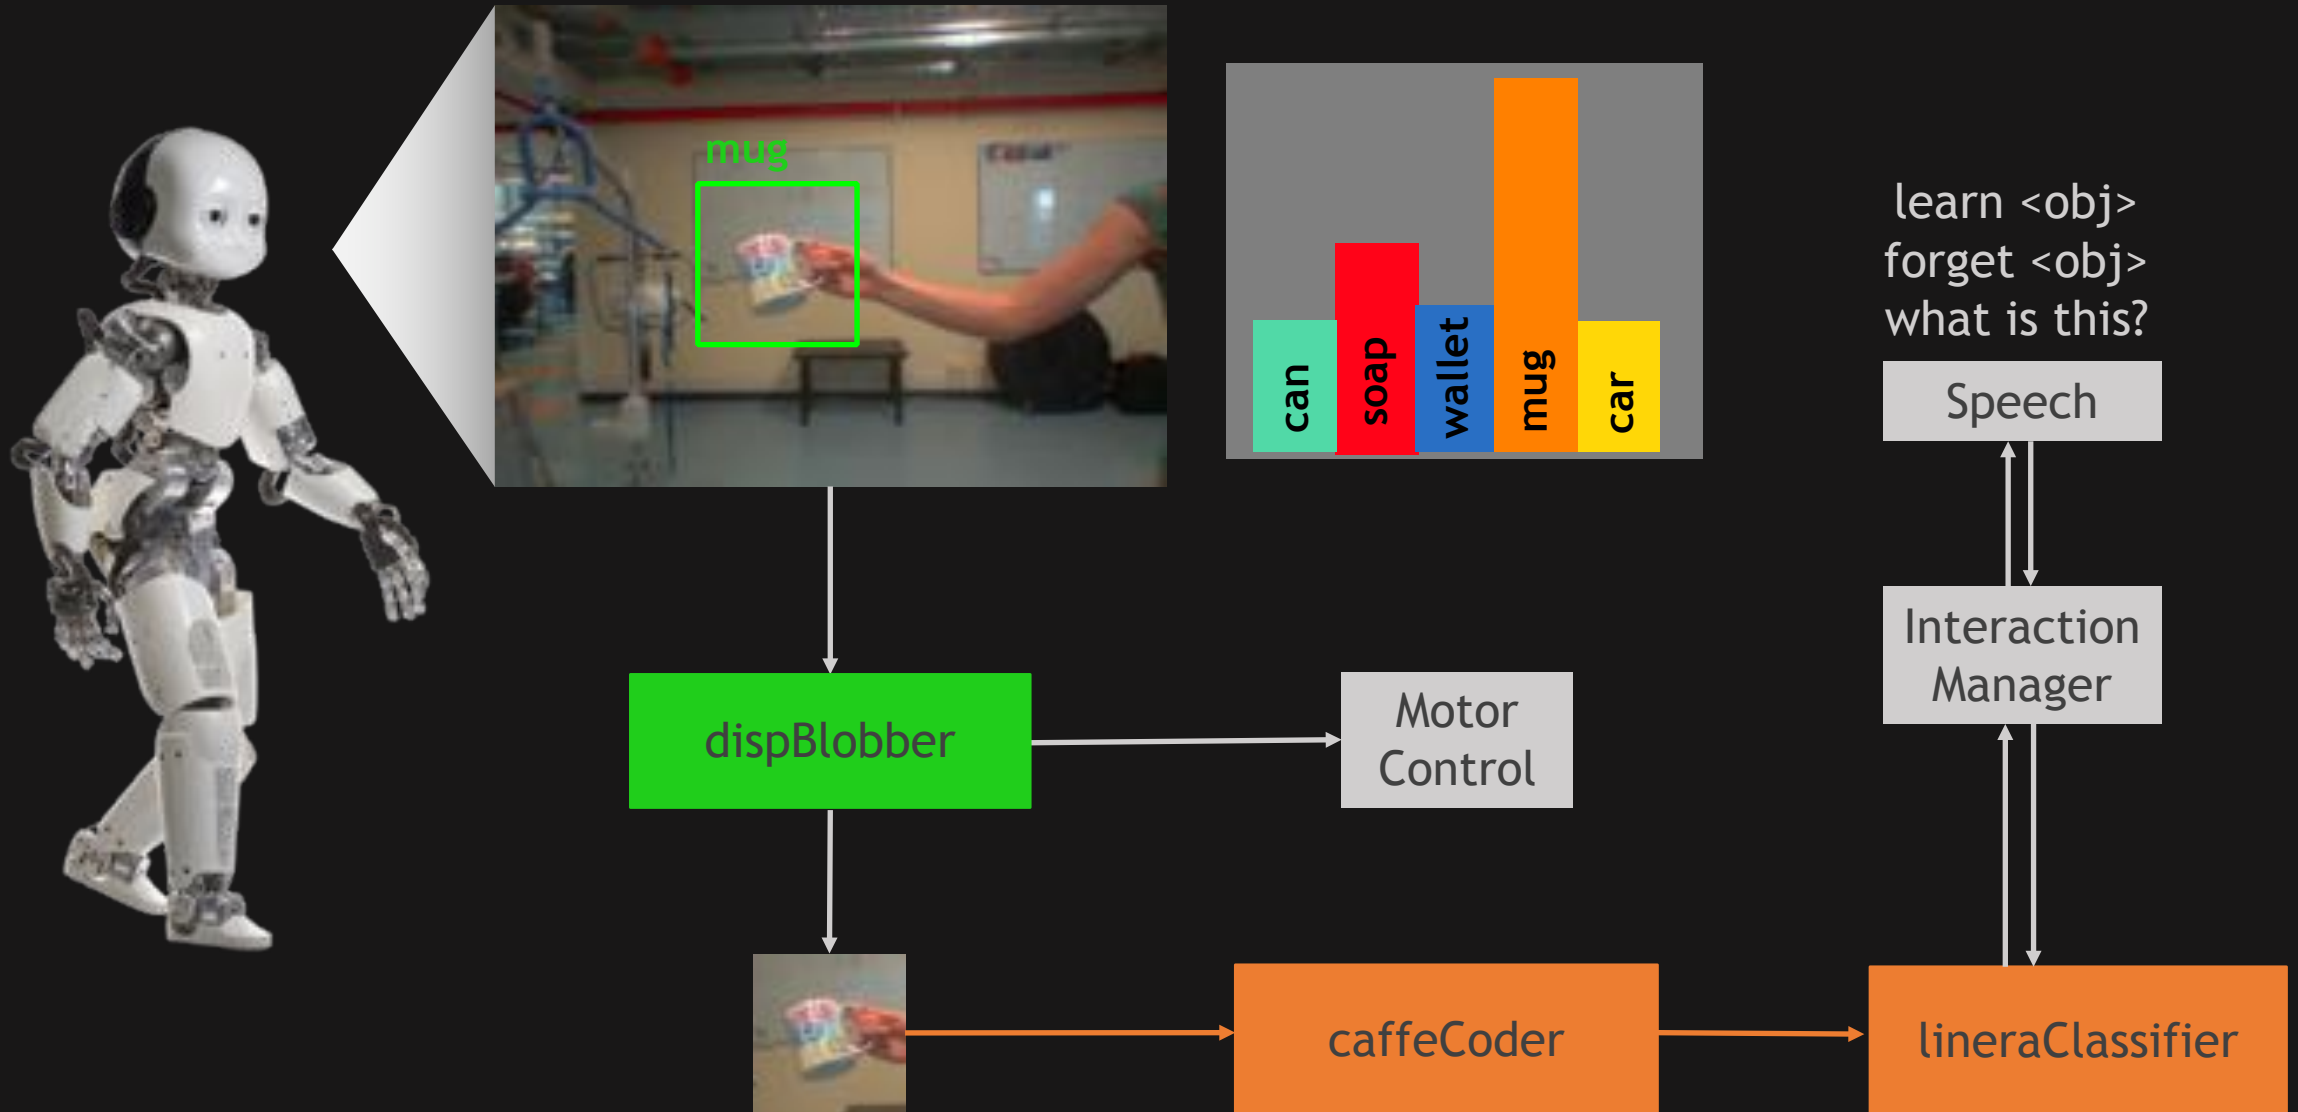
1. create a new exercise folder (call it e.g. `cat_3categories`) by copying `id_10objects_caffenet`
 2. open `id_10objects.m` in the new folder (you may rename it to match the new task) and change it to generate a task of discrimination between the 3 categories in iCW (e.g. by training on 7 objects/cat, validating on 2, and testing on the remaining 1)
 3. open `example_task_generation.m` and modify the paths to
 - (i) use the configuration file at point 2)
 - (ii) create training, validation and test sets in the new folder→ run the script to generate the new training set
 4. implement the protocols `all-0` and `all-3` on this new task and run them
- how much does the recognition accuracy change? is there a better protocol?
- is this task more or less difficult? why?

Things you may want to try...

Deploy the fine-tuned model in the recognition pipeline on the iCub!

- create a folder with your surname
- put inside the `final.caffemodel` and `mean.binaryproto`
- together with a `test.prototxt` (see next for how to create this file...)

On the Fly Object Recognition Demo on iCub



Pointers to GitHub Repositories

DEMO: [robotology/onthefly-recognition](https://github.com/robotology/onthefly-recognition)

- manager
- speech
- tracking


SEGMENTATION: [robotology/segmentation](https://github.com/robotology/segmentation)

- dispBlobber

RECOGNITION PIPELINE: [robotology/himrep](https://github.com/robotology/himrep)

- caffeCoder
- linearClassifier

caffeCoder: (YARP) wrapper for caffe



caffeCoder: call with your own files

provided that you have:

```
/home/icub/.local/share/yarp/contexts/himrep/pasquale/final.caffemodel  
/home/icub/.local/share/yarp/contexts/himrep/pasquale/mean.binaryproto  
/home/icub/.local/share/yarp/contexts/himrep/pasquale/test.prototxt
```


call (you can try also on the VM):

```
caffeCoder --pretrained_binary_proto_file  
/home/icub/.local/share/yarp/contexts/himrep/pasquale/final.caffemodel  
--feature_extraction_proto_file pasquale/test.prototxt --extract_features_blob_names fc6N
```


More details on assignment...

1. Fine-tuning CaffeNet on a ~~2~~-class identification task

10

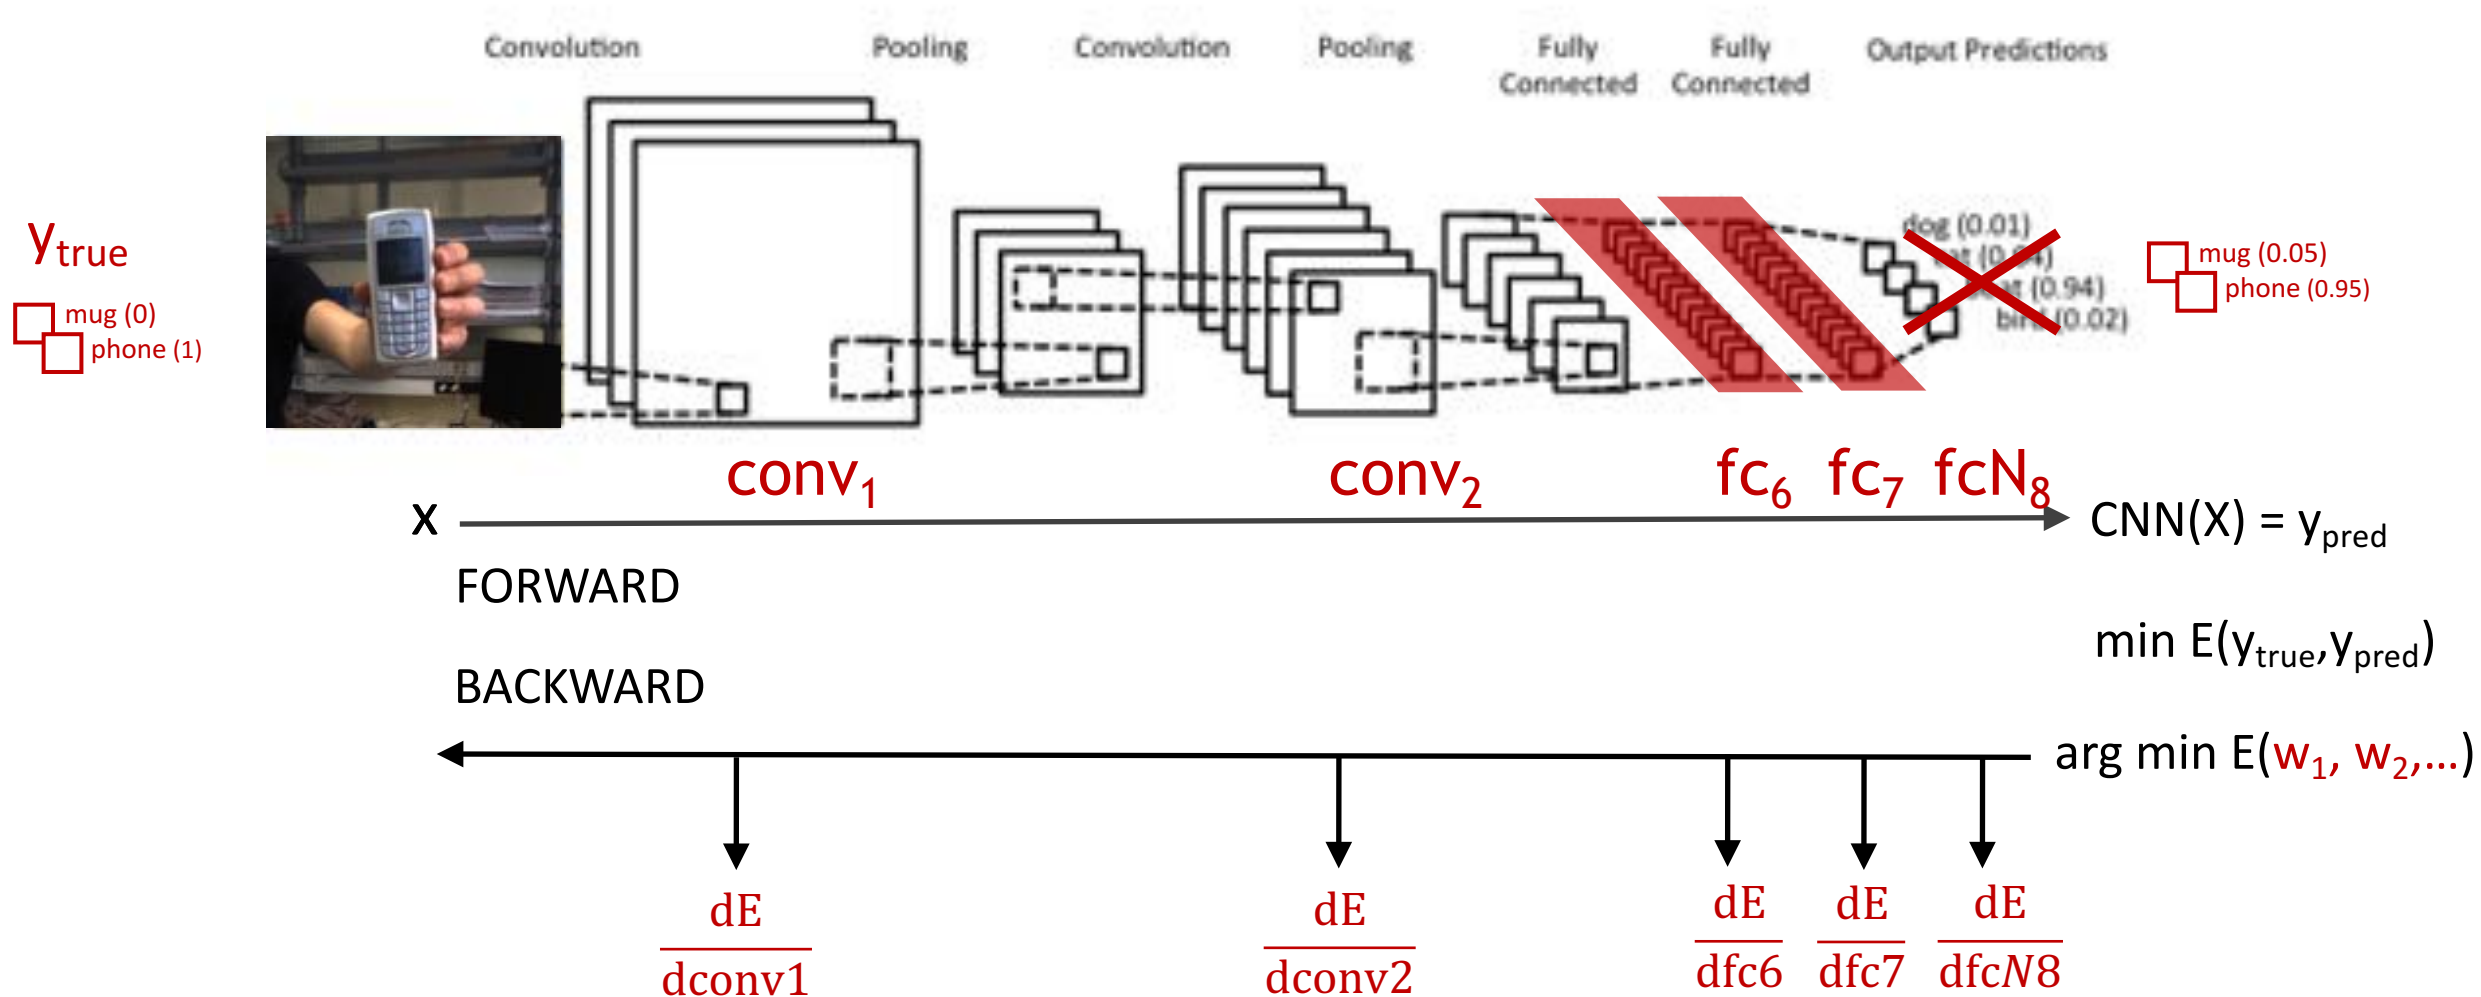


1. Fine-tuning CaffeNet on a ~~2~~-class identification task 10




name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
all-0	32	lr: 0	lr: 0	lr: 0	lr: 1e-2 from scratch	6	?	?

2. Try fine-tuning fc7,fc6

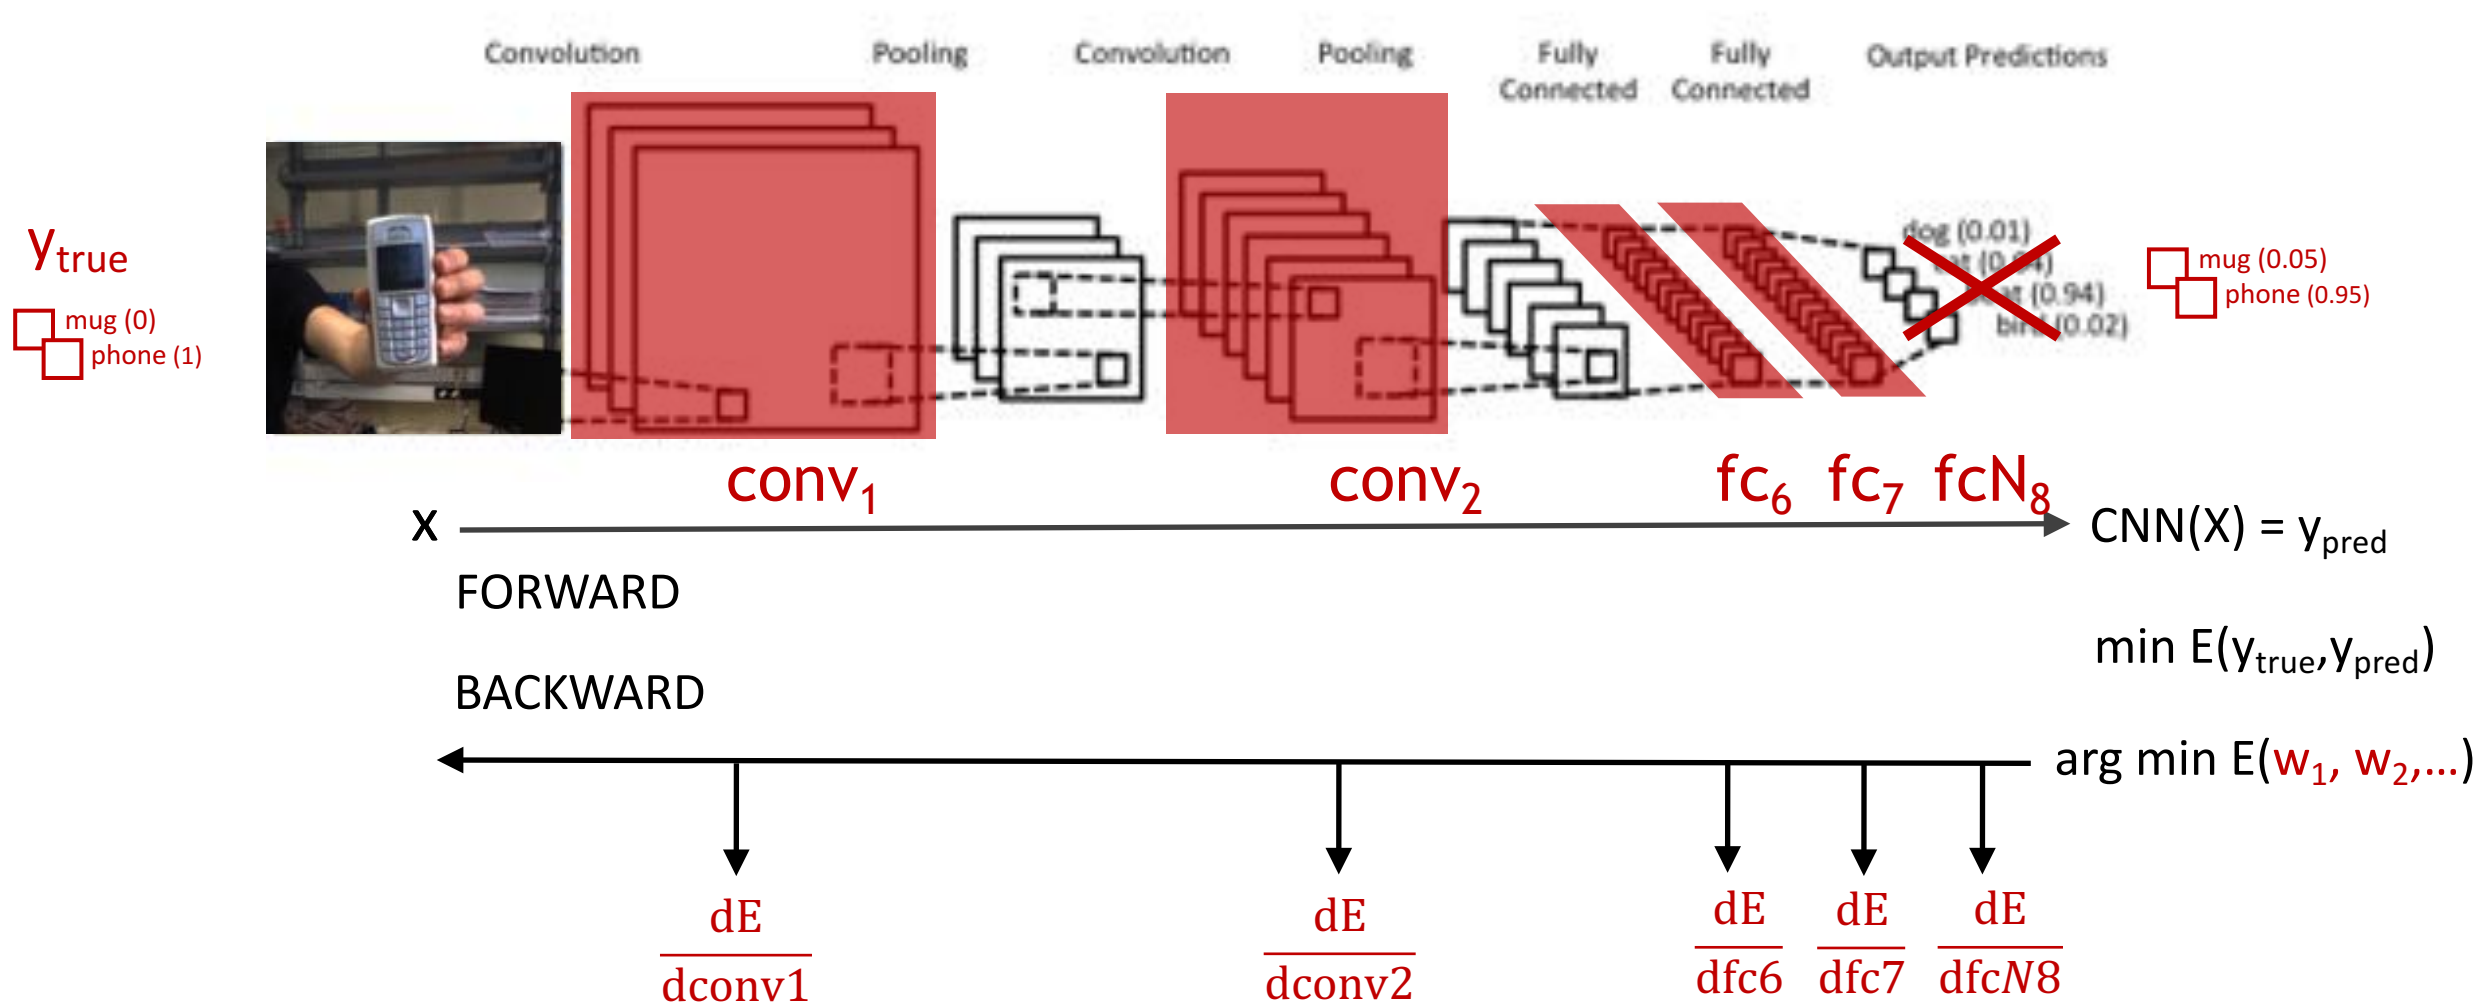


2. Try fine-tuning fc7,fc6

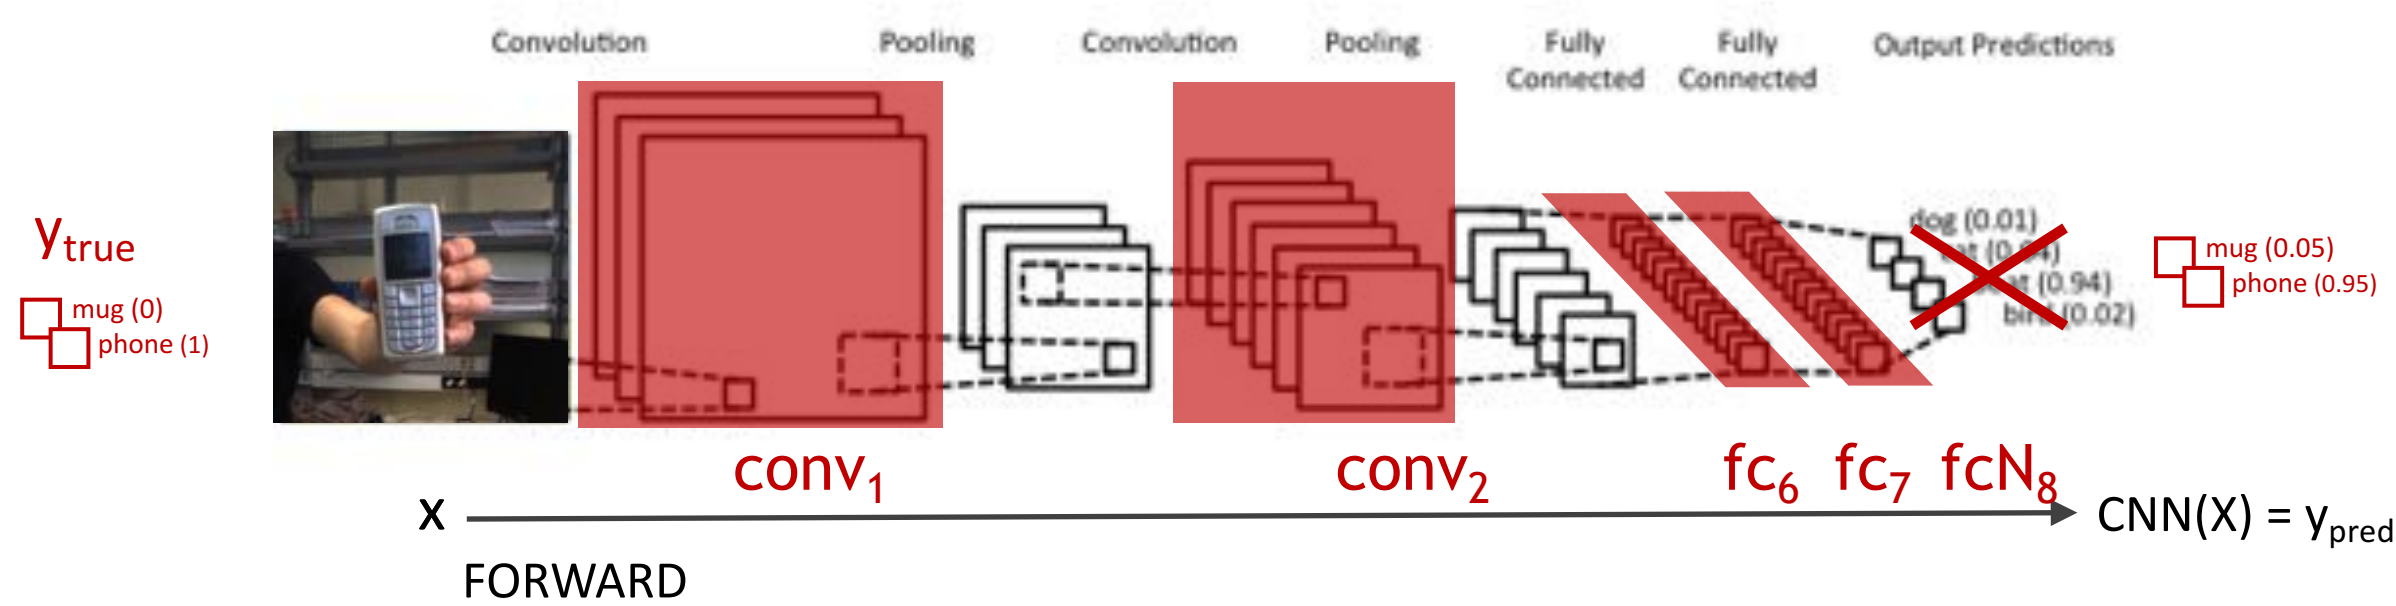


name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
conv-0_fc6-2_fc7-2	32	lr: 0	lr: 1e-2	lr: 1e-2	lr: 1e-2 from scratch	6	?	?
conv-0_fc6-3_fc7-3	32	lr: 0	lr: 1e-3	lr: 1e-3	lr: 1e-2 from scratch	6	?	?

3. Try fine-tuning all layers

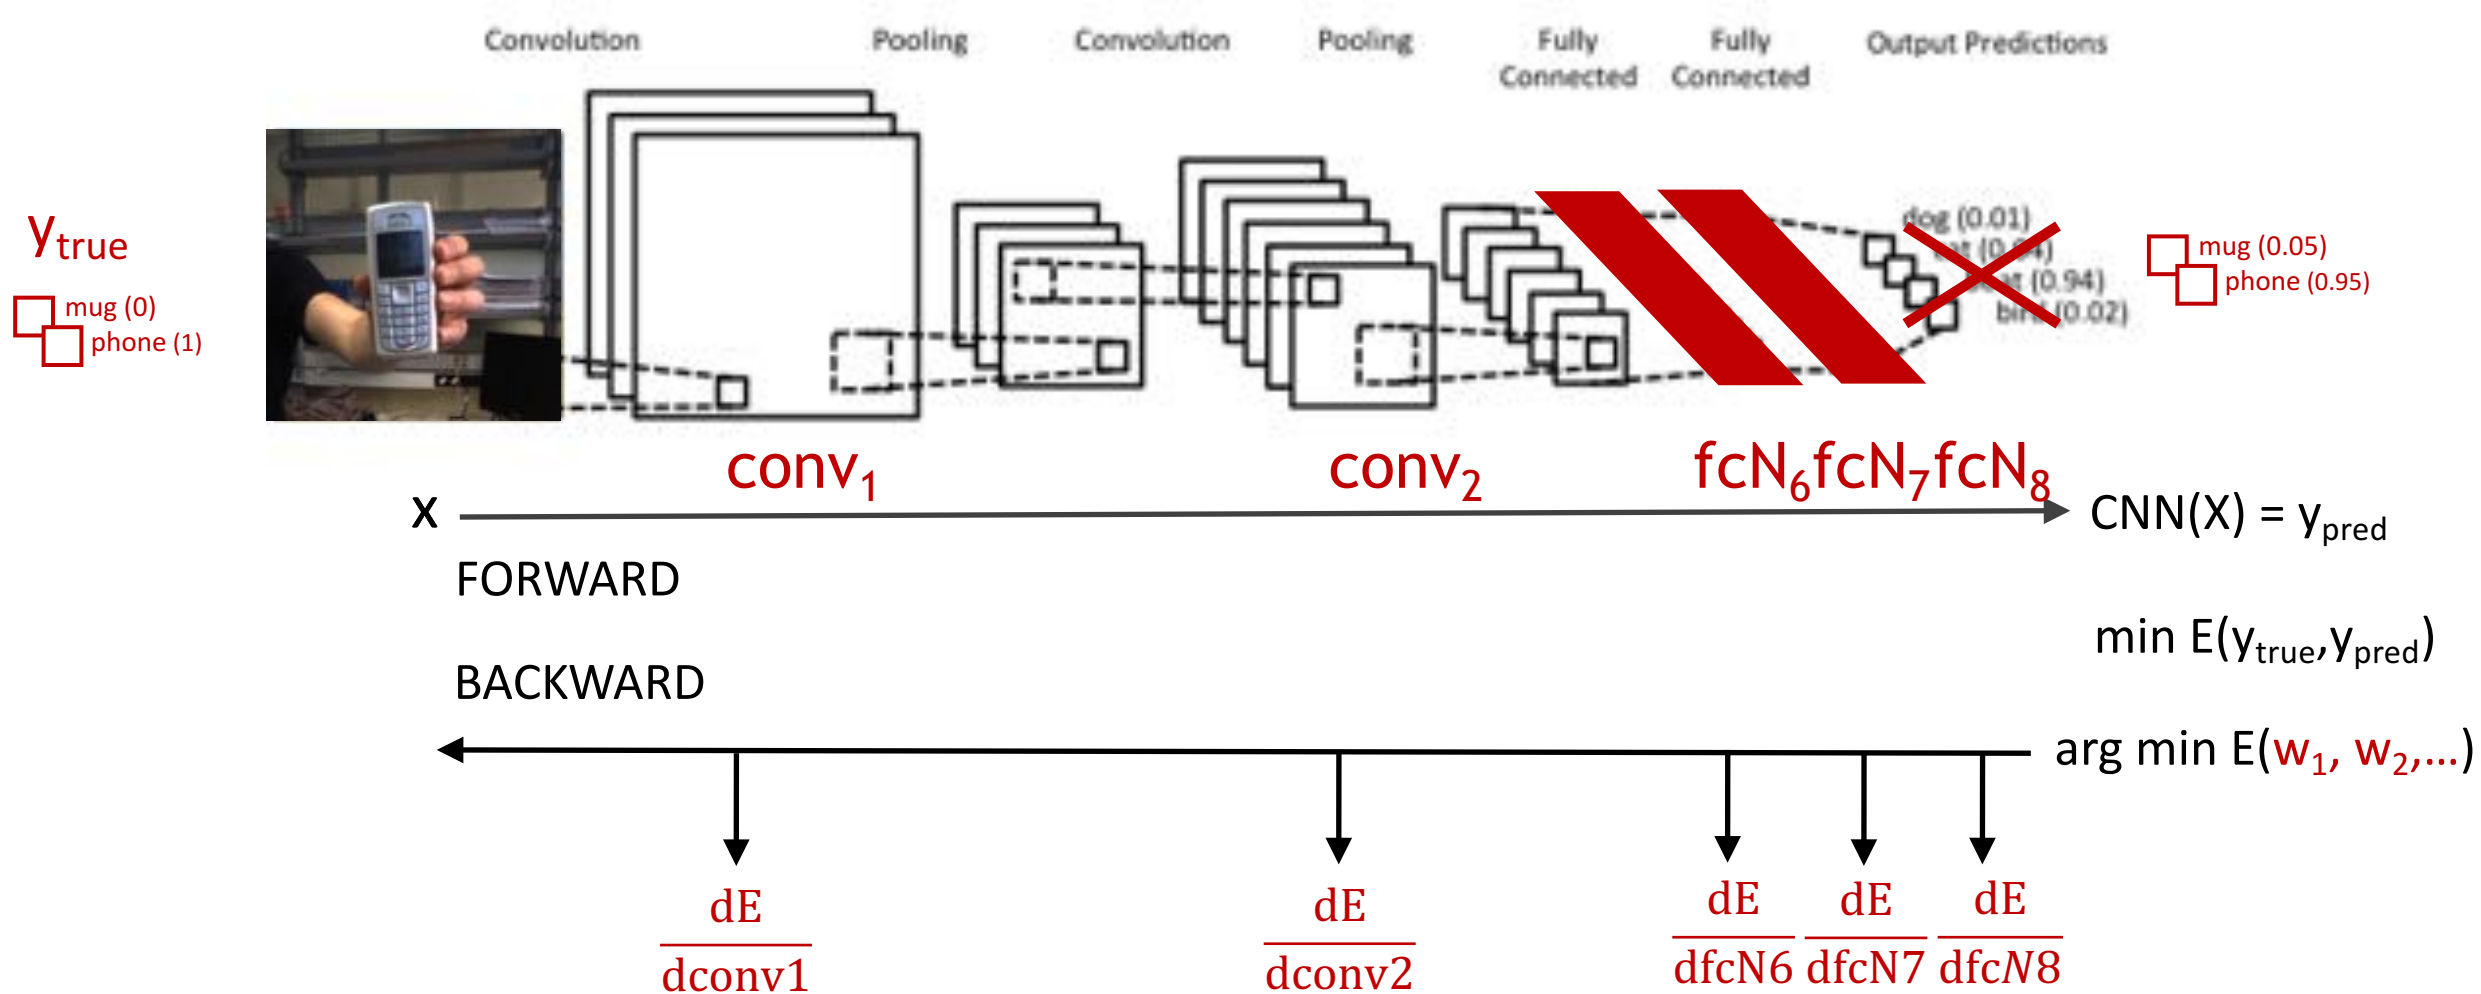


3. Try fine-tuning all layers

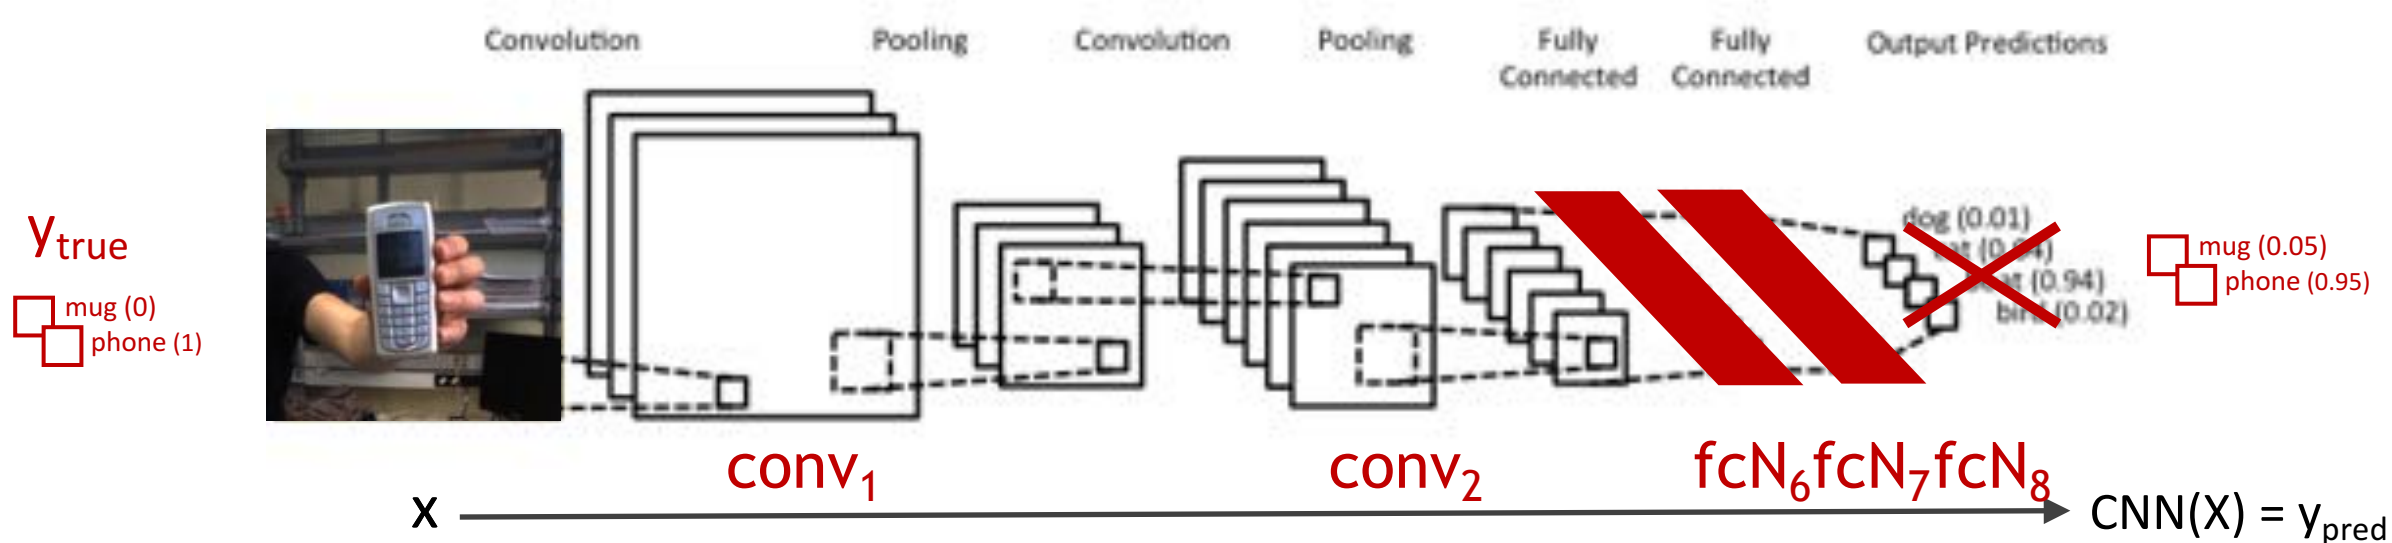


name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
all-3	32	lr: 1e-3	lr: 1e-3	lr: 1e-3	lr: 1e-2 from scratch	6 (12)	?	?

4. Is it possible to learn from scratch $fc7, fc6$?



4. Is it possible to learn from scratch fc7,fc6?



name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
conv-0_fc6N-2_fc7N-2	32	lr: 0	lr: 1e-2 from scratch	lr: 1e-2 from scratch	lr: 1e-2 from scratch	18	?	?
conv-0_fc6N-4_fc7N-4	32	lr: 0	lr: 1e-4 from scratch	lr: 1e-4 from scratch	lr: 1e-2 from scratch	18 (36)	?	?

5. Vary the batch size

name	batch size (train)	conv layers	fc6	fc7	fc8	# epochs	acc day1	acc day2
all-3_batch-8	8	lr: 1e-3	lr: 1e-3	lr: 1e-3	lr: 1e-2 from scratch	6	?	?
all-0_batch-1	1	lr: 0	lr: 0	lr: 0	lr: 1e-2 from scratch	6	?	?

`solver.prototxt`: how to set batch-related parameters

```
net: "train_val.prototxt"
```

```
solver_mode: CPU
```

```
# carry out <max_iter> training iterations
```

```
max_iter: 144 (6 x 24)
```

```
## the validation will carry out <test_iter> iterations
```

```
test_iter: 10
```

```
## carry out validation every <test_interval> training iterations
```

```
test_interval: 24
```

```
## display every <display> iterations
```


```
display: 24
```

```
## save model every <snapshot> iterations
```

```
snapshot: 0
```

```
snapshot_prefix: "icw"
```

batch_size



(1 epoch = 24 iterations)

training set:

$n = 768$

$\text{batch_size} = 32$

$\text{iters_per_epoch} = 768/32 = 24$

$\text{max_epoch} = 6$

validation set:

$n = 320$

$\text{batch_size} = 32$

$\text{iters_per_epoch} = 320/32 = 10$

$\text{test_epoch} = 1$

train_val.prototxt: set the batch size

```
layer {
  name: "data"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TRAIN
  }
  transform_param {
    mirror: true
    crop_size: 227
    mean_file: "../mean.binaryproto"
  }
  data_param {
    source: "../lmdb_train"
    batch_size: 32
    backend: LMDB
  }
}
```

```
layer {
  name: "data"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TEST
  }
  transform_param {
    mirror: false
    crop_size: 227
    mean_file: "../mean.binaryproto"
  }
  data_param {
    source: "../lmdb_val"
    batch_size: 32
    backend: LMDB
  }
}
```


train_val.prototxt: set lr_mult for conv and fc layers

```
layer {  
  name: "conv1"  
  type: "Convolution"  
  bottom: "data"  
  top: "conv1"  
  param {  
    lr_mult: <x> [1,10,0.1...]  
    decay_mult: 1  
  }  
  param {  
    lr_mult: <2x> [2,20,0.2...]  
    decay_mult: 0  
  }  
  ... ..
```

```
... ..  
  convolution_param {  
    num_output: 96  
    kernel_size: 11  
    stride: 4  
    weight_filler {  
      type: "gaussian"  
      std: 0.01  
    }  
    bias_filler {  
      type: "constant"  
      value: 0  
    }  
  }  
}
```

train_val.prototxt: set lr_mult for conv and fc layers

```
layer {
  name: "fc6"
  type: "InnerProduct"
  bottom: "pool5"
  top: "fc6"
  param {
    lr_mult: <x> [e.g.: 1,10,0.1...]
    decay_mult: 1
  }
  param {
    lr_mult: <2x> [e.g.: 2,20,0.2...]
    decay_mult: 0
  }
  ... ..
```

```
... ..
  inner_product_param {
    num_output: 4096
    weight_filler {
      type: "gaussian"
      std: 0.005
    }
    bias_filler {
      type: "constant"
      value: 1
    }
  }
}
```

train_val.prototxt: rename layers (and related output blobs)
that you learn from scratch

```
layer {  
  name: "fc8N"  
  type: "InnerProduct"  
  bottom: "fc7"  
  top: "fc8N"  
  param {  
    lr_mult: 1  
    decay_mult: 1  
  }  
  param {  
    lr_mult: 2  
    decay_mult: 0  
  }  
  ... ..
```

```
... ..  
  inner_product_param {  
    num_output: 2  
    weight_filler {  
      type: "gaussian"  
      std: 0.01  
    }  
    bias_filler {  
      type: "constant"  
      value: 0  
    }  
  }  
}
```

train_val.prototxt: rename layers (and related output blobs)
that you learn from scratch

```
layer {  
  name: "accuracy"  
  type: "Accuracy"  
  bottom: "fc8N"  
  bottom: "label"  
  top: "accuracy"  
}
```

```
layer {  
  name: "loss"  
  type: "SoftmaxWithLoss"  
  bottom: "fc8N"  
  bottom: "label"  
  top: "loss"  
}
```

deploy.prototxt

- ✓ same structure as `train_val.prototxt`, but

 - no DATA layer

 - no learning parameters (`lr_mult...`)

- ✓ pay attention to:

 - rename the layers which you learned from scratch (`fc8` → `fc8N`)

 - set `num_output (fc8N)` accordingly to the task (`num_output=2`)

test.prototxt: you have to create it

- ✓ same structure as `train_val.prototxt`, but

“MemoryData” DATA layer (see next slide)
no learning parameters

- ✓ pay attention to:

→ rename the layers which you learned from scratch (`fc8` → `fc8N`)

→ set `num_output` (`fc8N`) accordingly to the task (`num_output=2`)

(you can copy the `deploy.prototxt`)

test.prototxt: example “MemoryData” layer

```
layer {  
  name: "data"  
  type: "MemoryData"  
  top: "data"  
  top: "label"  
  transform_param {  
    mirror: false  
    crop_size: 227  
    mean_file: "/home/icub/.local/share/yarp/contexts/himrep/pasquale/mean.binaryproto "  
  }  
  memory_data_param {  
    batch_size: 1  
    channels: 3  
    height: 227  
    width: 227  
  }  
}
```